

1- Logical Volume Management

- a- Create a volume group and logical volume. Format it, mount it, and extend it by 1GB.

>> At first you must add a volumes from the vmware

```
[root@localhost direct]# lsblk -p
NAME                                MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
/dev/sr0                             11:0    1 116.9M  0 rom  /run/media/ibrahim/CDROM
/dev/sr1                             11:1    1  10.3G  0 rom  /run/media/ibrahim/RHEL-9-4-0-BaseOS-x86_64
/dev/nvme0n1                         259:0    0   30G   0 disk
├─/dev/nvme0n1p1                     259:1    0   600M   0 part /boot/efi
├─/dev/nvme0n1p2                     259:2    0    1G   0 part /boot
├─/dev/nvme0n1p3                     259:3    0  28.4G   0 part
│ └─/dev/mapper/rhel-root            253:0    0  25.4G   0 lvm  /
│   └─/dev/mapper/rhel-swap          253:1    0    3G   0 lvm  [SWAP]
└─/dev/nvme0n2                       259:4    0    5G   0 disk
/dev/nvme0n3                         259:5    0    5G   0 disk
/dev/nvme0n4                         259:6    0    5G   0 disk
/dev/nvme0n5                         259:7    0    5G   0 disk
[root@localhost direct]#
```

>>Then you need to label these devices as a pv

```
[root@localhost direct]# pvcreate /dev/nvme0n2 /dev/nvme0n3
Physical volume "/dev/nvme0n2" successfully created.
Physical volume "/dev/nvme0n3" successfully created.
[root@localhost direct]# pvs
PV          VG      Fmt  Attr PSize  PFree
/dev/nvme0n1p3 rhel  lvm2 a-- 28.41g  0
/dev/nvme0n2   lvm2 --- 5.00g 5.00g
/dev/nvme0n3   lvm2 --- 5.00g 5.00g
[root@localhost direct]#
```

>> Add these devices to a volume group

```
[root@localhost direct]# vgcreate mygroup /dev/nvme0n2 /dev/nvme0n3
Volume group "mygroup" successfully created
[root@localhost direct]# vgs
VG      #PV #LV #SN Attr   VSize  VFree
mygroup 2   0   0 wz--n- 9.99g 9.99g
rhel    1   2   0 wz--n- 28.41g  0
[root@localhost direct]#
```

>> Here each device is a 5GB so the group of these devices is about 10GB

>>Create a logical volume from the available free space on the volume group

```
[root@localhost direct]# lvcreate -n lv01 -L 3GB mygroup
WARNING: vdo signature detected on /dev/mygroup/lv01 at offset 0. Wipe it? [y/n]: y
Wiping vdo signature on /dev/mygroup/lv01.
Logical volume "lv01" created.
[root@localhost direct]# lvs
LV VG      Attr   LSize  Pool Origin Data%  Meta%  Move Log Cpy%Sync Convert
lv01 mygroup -wi-a---- 3.00g
root rhel -wi-ao---- 25.41g
swap rhel -wi-ao---- 3.00g
[root@localhost direct]#
```

>> First get the path of the logical volume and Choose a file system (xfs) to be assigned to this logical volume's path

```
Nov 3 21:15
ibrahim@localhost:shares/direct

[root@localhost direct]# lvsdisplay
--- Logical volume ---
LV Path                /dev/mygroup/lv01

Nov 3 21:17
ibrahim@localhost:~

[root@localhost ~]# mkfs -t xfs /dev/mygroup/lv01
meta-data=/dev/mygroup/lv01      isize=512    agcount=4, agsize=196608 blks
=                               sectsz=512   attr=2, projid32bit=1
=                               crc=1       finobt=1, sparse=1, rmapbt=0
=                               reflink=1   bigtime=1 inobtcount=1 nrext64=0
data      =                       bsize=4096   blocks=786432, imaxpct=25
=                               sunit=0      swidth=0 blks
naming    =version 2             bsize=4096   ascii-ci=0, ftype=1
log       =internal log         bsize=4096   blocks=16384, version=2
=                               sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none                  extsz=4096   blocks=0, rtextents=0
[root@localhost ~]#
```

>> Now you can put this file system and its logical volume in the fstab file to mount this logical volume and its file system on a mount point (lv01_data) and make it permanent across system reboots

```
Nov 3 21:22
ibrahim@localhost:~

[root@localhost ~]# mkdir /lv01_data

Nov 3 21:25
ibrahim@localhost:~

[root@localhost ~]# nano /etc/fstab

GNU nano 5.6.1 /etc/fstab Modified

#
# /etc/fstab
# Created by anaconda on Mon Oct 28 18:33:01 2024
#
# Accessible filesystems, by reference, are maintained under '/dev/disk/'.
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info.
#
# After editing this file, run 'systemctl daemon-reload' to update systemd
# units generated from this file.
#
/dev/mapper/rhel-root / xfs defaults 0 0
UUID=0f570de3-b3e0-40a2-bf3f-76c1610dd7ad /boot xfs defaults 0 0
UUID=8D9C-2D77 /boot/efi vfat umask=0077,shortname=winnt 0 2
/dev/mapper/rhel-swap none swap defaults 0 0
/dev/mygroup/lv01 /lv01_data xfs defaults 0 0
```

>> Save and exit

```
Nov 3 21:27
ibrahim@localhost:~

[root@localhost ~]# mount -a
mount: (hint) your fstab has been modified, but systemd still uses
the old version; use 'systemctl daemon-reload' to reload.
[root@localhost ~]# systemctl daemon-reload
[root@localhost ~]# df -h
Filesystem                Size      Used Avail Use% Mounted on
devtmpfs                   4.0M         0   4.0M   0% /dev
tmpfs                      1.9G         0   1.9G   0% /dev/shm
tmpfs                      777M       9.8M   767M   2% /run
/dev/mapper/rhel-root      26G       6.8G    19G  27% /
/dev/nvme0n1p2             960M      264M   697M  28% /boot
/dev/nvme0n1p1             599M       7.1M   592M   2% /boot/efi
tmpfs                      1.0M         0    1.0M   0% /run/stratisd/ns_mounts
tmpfs                      389M      100K   389M   1% /run/user/1000
/dev/sr0                   117M       117M    0 100% /run/media/ibrahim/CDROM
/dev/sr1                   11G        11G    0 100% /run/media/ibrahim/RHEL-9-4-0-BaseOS-x86_64
/dev/mapper/mygroup-lv01  3.0G       54M    2.9G   2% /lv01_data
[root@localhost ~]#
```

>>To extend this logical volume by 1GB

```
ibrahim@localhost:~$ sudo lvextend -r -L +1G /dev/mygroup/lv01
Size of logical volume mygroup/lv01 changed from 3.00 GiB (768 extents) to 4.00 GiB (1024 extents).
File system xfs found on mygroup/lv01 mounted at /lv01_data.
Extending file system xfs to 4.00 GiB (4294967296 bytes) on mygroup/lv01...
xfs_growfs /dev/mygroup/lv01
meta-data=/dev/mapper/mygroup-lv01 isize=512    agcount=4, agsize=196608 blks
       =                               sectsz=512   attr=2, projid32bit=1
       =                               crc=1        finobt=1, sparse=1, rmapbt=0
       =                               reflink=1    bigtime=1 inobtcount=1 nrext64=0
data      =                               bsize=4096 blocks=786432, imaxpct=25
       =                               sunit=0      swidth=0 blks
naming    =version 2                       bsize=4096   ascii-ci=0, ftype=1
log       =internal log                   bsize=4096   blocks=16384, version=2
       =                               sectsz=512   sunit=0 blks, lazy-count=1
realtime  =none                           extsz=4096    blocks=0, rtextents=0
data blocks changed from 786432 to 1048576
xfs_growfs done
Extended file system xfs on mygroup/lv01.
Logical volume mygroup/lv01 successfully resized.
ibrahim@localhost:~$ df -h
Filesystem                Size      Used Avail Use% Mounted on
devtmpfs                  4.0M        0  4.0M   0% /dev
tmpfs                     1.9G        0  1.9G   0% /dev/shm
tmpfs                     777M    9.8M  767M   2% /run
/dev/mapper/rhel-root      26G    6.8G   19G  27% /
/dev/nvme0n1p2            960M   264M  697M  28% /boot
/dev/nvme0n1p1            599M    7.1M  592M   2% /boot/efi
tmpfs                     1.0M        0   1.0M   0% /run/stratisd/ns_mounts
tmpfs                     389M   100K  389M   1% /run/user/1000
/dev/sr0                  117M   117M    0 100% /run/media/ibrahim/CDROM
/dev/sr1                   11G     11G    0 100% /run/media/ibrahim/RHEL-9-4-0-BaseOS-x86_64
/dev/mapper/mygroup-lv01  4.0G     61M   3.9G   2% /lv01_data
ibrahim@localhost:~$
```

>> Here is the verification about the storage used and the free storage in the vg

```
ibrahim@localhost:~$ sudo vgdisplay
--- Volume group ---
VG Name                mygroup
System ID
Format                 lvm2
Metadata Areas         2
Metadata Sequence No   3
VG Access               read/write
VG Status               resizable
MAX LV                 0
Cur LV                 1
Open LV                 1
Max PV                 0
Cur PV                 2
Act PV                  2
VG Size                9.99 GiB
PE Size                 4.00 MiB
Total PE                2558
Alloc PE / Size         1024 / 4.00 GiB
Free PE / Size          1534 / 5.99 GiB
VG UUID                nF1TFF-kDYS-01Z1-NuT1-khIW-vWWZ-oi0ZdN
```

2- NFS Setup

- a- Install and configure an NFS server and client. Share a directory and verify accessibility from the client machine.

>> In order to achieve this requirement I have a two vms up and running
Ubuntu as a NFS client

RHEL 9.4 as NFS server

>> The common step must be executed on the two vm is installing the NFS package.

>> First let's setup the server side

>> install and enable the nfs-server

```
Nov 3 22:14
ibrahim@localhost:~
ibrahim@localhost:~
ibrahim@localhost:~

[root@localhost ~]# dnf install -y nfs-utils
Updating Subscription Management repositories.

This system is registered with an entitlement server, but is not receiving updates. You can use subscription-man
ager to assign subscriptions.

Last metadata expiration check: 0:00:46 ago on Sun 03 Nov 2024 10:12:58 PM EET.
Package nfs-utils-1:2.5.4-26.el9_4.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
[root@localhost ~]# systemctl enable --now nfs-server
[root@localhost ~]# systemctl status nfs-server
● nfs-server.service - NFS server and services
   Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; preset: disabled)
   Drop-In: /run/systemd/generator/nfs-server.service.d
            └─order-with-mounts.conf
   Active: active (exited) since Sun 2024-11-03 19:36:09 EET; 2h 38min ago
     Docs: man:rpc.nfsd(8)
           man:exportfs(8)
    Main PID: 1294 (code=exited, status=0/SUCCESS)
      CPU: 24ms

Nov 03 19:36:08 localhost.localdomain systemd[1]: Starting NFS server and services...
Nov 03 19:36:09 localhost.localdomain systemd[1]: Finished NFS server and services.
[root@localhost ~]#
```

>> Enable the nfs-servie on the firewall

```
Nov 3 22:19
ibrahim@localhost:~
ibrahim@localhost:~
ibrahim@localhost:~

[root@localhost ~]# firewall-cmd --permanent --add-service=nfs
Warning: ALREADY_ENABLED: nfs
success
[root@localhost ~]# firewall
firewall-cmd      firewallld      firewall-offline-cmd
[root@localhost ~]# firewall-cmd --list-all
public (active)
target: default
icmp-block-inversion: no
interfaces: ens160
sources:
services: cockpit dhcpv6-client nfs ssh
ports:
protocols:
forward: yes
masquerade: no
forward-ports:
source-ports:
icmp-blocks:
rich rules:
[root@localhost ~]#
```

>>Create the shared directory

```
ibrahim@localhost:~$ sudo mkdir /shared_data
ibrahim@localhost:~$ ls -l /shared_data/
total 0
ibrahim@localhost:~$ ls -ld /shared_data
drwxr-xr-x. 2 root root 6 Nov  3 22:21 /shared_data
ibrahim@localhost:~$ # give others the write permission to allow client create files on the shared directory
ibrahim@localhost:~$ sudo chmod o+w /shared_data
ibrahim@localhost:~$ ls -ld /shared_data/
drwxr-xrwx. 2 root root 6 Nov  3 22:21 /shared_data/
ibrahim@localhost:~$
```

>>Get the client IP and Export the shared directory

```
ibrahim@ibrahim-server:~$ nmcli device show
GENERAL.DEVICE: ens33
GENERAL.TYPE: ethernet
GENERAL.HWADDR: 00:0C:29:A6:BE:C8
GENERAL.MTU: 1500
GENERAL.STATE: 100 (connected)
GENERAL.CONNECTION: netplan-ens33
GENERAL.CON-PATH: /org/freedesktop/NetworkManager/ActiveConnection/6
WIRED-PROPERTIES.CARRIER: on
IP4.ADDRESS[1]: 192.168.115.128/24
IP4.GATEWAY: 192.168.115.2
IP4.ROUTE[1]: dst = 192.168.115.0/24, nh = 0.0.0.0, mt = 100
IP4.ROUTE[2]: dst = 0.0.0.0/0, nh = 192.168.115.2, mt = 100
IP4.DNS[1]: 192.168.115.2
IP4.DOMAIN[1]: localdomain
IP6.ADDRESS[1]: fe80::20c:29ff:fea6:bec8/64
IP6.GATEWAY: --
IP6.ROUTE[1]: dst = fe80::/64, nh = ::, mt = 256

ibrahim@localhost:~$ nano /etc/exports
GNU nano 5.6.1 /etc/exports Modified
/shared_data 192.168.115.128(rw)
```

>>save and exit

>> Now use the exportfs -r to make the server re read the configuration file without terminating any connections

```
ibrahim@localhost:~$ nano /etc/exports
ibrahim@localhost:~$ exportfs -r
ibrahim@localhost:~$ exportfs
/shared_data 192.168.115.128
ibrahim@localhost:~$
```

>> Now let's configure the client side

>> Install the NFS package

```
Nov 3 22:40
root@ibrahim-server: ~
root@ibrahim-server:~# apt install -y nfs-server
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Note, selecting 'nfs-kernel-server' instead of 'nfs-server'
nfs-kernel-server is already the newest version (1:2.6.4-3ubuntu5).
0 upgraded, 0 newly installed, 0 to remove and 91 not upgraded.
root@ibrahim-server:~# systemctl status nfs-server
● nfs-server.service - NFS server and services
   Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; preset: enabled)
   Active: active (exited) since Sun 2024-11-03 19:35:50 EET; 3h 4min ago
     Main PID: 1843 (code=exited, status=0/SUCCESS)
        CPU: 75ms
Nov 03 19:35:50 ibrahim-server systemd[1]: Starting nfs-server.service - NFS server and services...
Nov 03 19:35:50 ibrahim-server systemd[1]: Finished nfs-server.service - NFS server and services.
root@ibrahim-server:~#
```

>> Now get the IP address of the NFS-server and create a mount point

“/NFS_mounted_data” on the client then go to the fstab file and add the mount record to make it exist across system reboots

```
Nov 3 22:43
ibrahim@localhost: ~
ibrahim@localhost:~# ifconfig
ens160: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.115.129 netmask 255.255.255.0 broadcast 192.168.115.255
    inet6 fe80::20c:29ff:fe51:7d3b prefixlen 64 scopeid 0x20<link>
    ether 00:0c:29:51:7d:3b txqueuelen 1000 (Ethernet)
    RX packets 2408 bytes 304116 (296.9 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 990 bytes 138150 (134.9 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

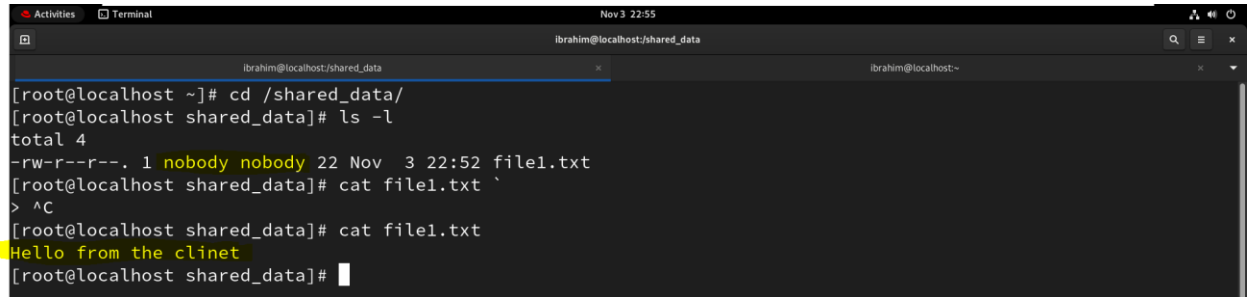
```
Nov 3 22:48
root@ibrahim-server: ~
GNU nano 7.2 /etc/fstab *
# /etc/fstab: static file system information.
#
# Use 'blkid' to print the universally unique identifier for a
# device; this may be used with UUID= as a more robust way to name devices
# that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point> <type> <options>          <dump> <pass>
# / was on /dev/sda2 during curtin installation
/dev/disk/by-uuid/37b2caa5-5e38-4f95-b187-e0579acbac50 / ext4 defaults 0 1
/swap.img none swap sw 0 0
192.168.115.129:/shared_data /NFS_mounted_data nfs defaults 0 0
```

```
root@ibrahim-server:~# mount -a
mount: (hint) your fstab has been modified, but systemd still uses
the old version; use 'systemctl daemon-reload' to reload.
root@ibrahim-server:~# systemctl daemon-reload
```

>> Verify that your client can reach the server's shared file

```
Nov 3 22:52
root@ibrahim-server: NFS_mounted_data
root@ibrahim-server:~# df -h
Filesystem      Size  Used Avail Use% Mounted on
tmpfs            387M  2.1M  385M   1% /run
/dev/sda2        20G   11G   7.7G  59% /
tmpfs            1.9G   0  1.9G   0% /dev/shm
tmpfs            5.0M   8.0K   5.0M   1% /run/lock
tmpfs            387M  148K  387M   1% /run/user/1000
192.168.115.129:/shared_data 26G   6.8G   19G  27% /NFS_mounted_data
root@ibrahim-server:~# cd /NFS_mounted_data/
root@ibrahim-server:/NFS_mounted_data# ls
root@ibrahim-server:/NFS_mounted_data# echo "Hello from the client" > file1.txt
root@ibrahim-server:/NFS_mounted_data#
```

>> Go to the server and check if the file that created on the client device is correctly created or not

A terminal window titled 'Terminal' with a timestamp of 'Nov 3 22:55'. The terminal shows a user 'ibrahim' at 'localhost' in the directory '/shared_data'. The user runs 'cd /shared_data/' and then 'ls -l', which shows a file 'file1.txt' owned by 'nobody' and created on '22 Nov 3 22:52'. The user then runs 'cat file1.txt', which outputs 'Hello from the client'.

```
[root@localhost ~]# cd /shared_data/
[root@localhost shared_data]# ls -l
total 4
-rw-r--r--. 1 nobody nobody 22 Nov  3 22:52 file1.txt
[root@localhost shared_data]# cat file1.txt
> ^C
[root@localhost shared_data]# cat file1.txt
Hello from the client
[root@localhost shared_data]#
```

>> The “nobody” user appears on the server is: Because of the client’s root user has created this file plus I have chosen the root squash option on the NFS-server for security reasons the client’s root user is assigned to “nobody” user which is a regular user on my NFS-server.